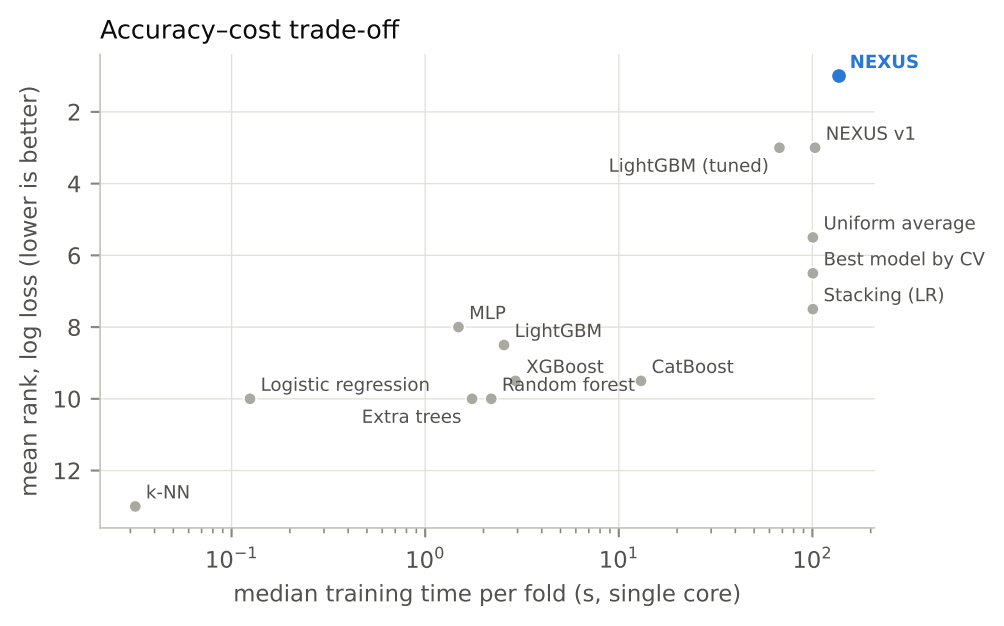
Accuracy–cost trade-off
NEXUS
NEXUS v1
LightGBM (tuned)
Uniform average
Best model by CV
Stacking (LR)
MLP
LightGBM
XGBoost
CatBoost
Logistic regression
Random forest
Extra trees
k-NN
mean rank, log loss (lower is better)
2
4
6
8
10
12
10^{-1}
10^{0}
10^{1}
10^{2}
median training time per fold (s, single core)