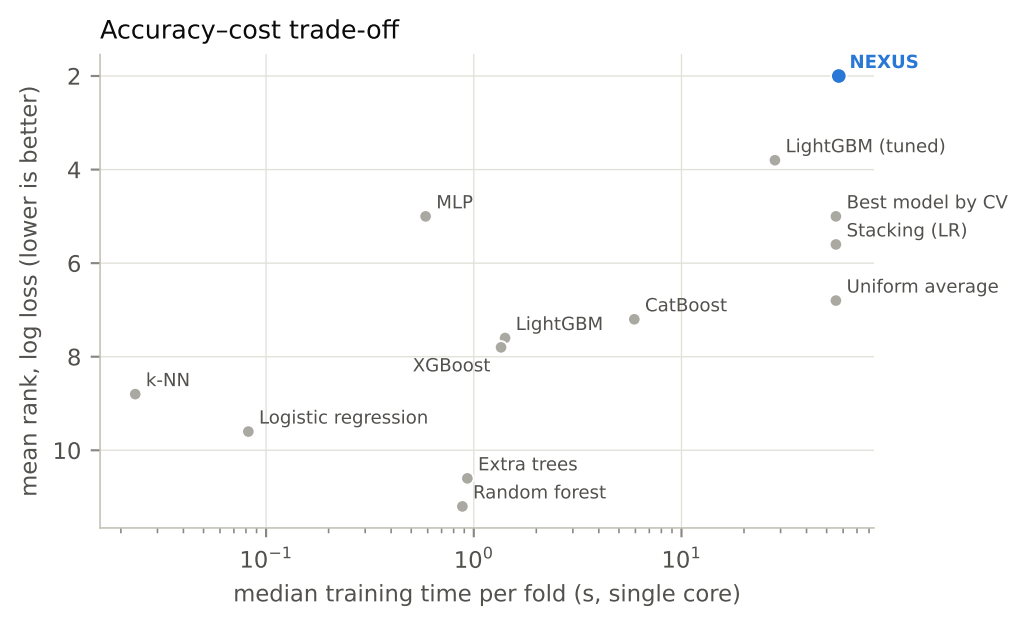
Accuracy–cost trade-off
NEXUS
LightGBM (tuned)
MLP
Best model by CV
Stacking (LR)
Uniform average
CatBoost
LightGBM
XGBoost
k-NN
Logistic regression
Extra trees
Random forest
2
4
6
8
10
mean rank, log loss (lower is better)
10^{-1}
10^{0}
10^{1}
median training time per fold (s, single core)